

EXECUTIVE SUMMARY

AI Behavioral Safety — Board Brief

A condensed snapshot of risk posture and remediation impact from a completed Ikwe audit. System names and sensitive details are redacted. The structure and findings are intact and tied to benchmark methodology.

BASELINE RISK	AFTER MITIGATION	REPAIR GAP
54.7% passed Safety Gate · 45.3% introduced harm Baseline measured across n=948 responses in 79 scenarios, pre-mitigation.	42–67% post-control reduction range Observed across scored risk dimensions.	43% no-repair-after-harm rate Baseline share of harmful responses with no repair action.

RISK POSTURE SCORECARD



Baseline vs post-mitigation scorecard across five risk dimensions. Safety Gate is a separate binary control check.

RISK DIMENSIONS

- Emotional Escalation — elevated before controls
- Dependency Formation — reduced after mitigation
- Authority Drift — stabilized with governance
- Scale Amplification — moderated under safeguards
- Governance Failure — low after controls