

The Behavioral Safety Layer for Human-Facing AI

Third-party independent behavioral safety validation for human-facing AI systems.

Benchmark scope: 79 scenarios | 13 behavioral domains (vulnerability categories).

Send completed PDF to research@ikwe.ai or submit online at ikwe.ai/intake.

A. Contact + Organization

Full name

Role / title

Company

Work email

Technical contact name

Technical contact email

Company size

Industry

Country / region

The Behavioral Safety Layer for Human-Facing AI

Third-party independent behavioral safety validation for human-facing AI systems.

Benchmark scope: 79 scenarios | 13 behavioral domains (vulnerability categories).

Send completed PDF to research@ikwe.ai or submit online at ikwe.ai/intake.

B. Deployment Context

Primary use case

What are you deploying?

Is it user-facing?

User population (select all that apply)

General consumers

Patients

Employees

Students

Other user population

Where is it deployed?

The Behavioral Safety Layer for Human-Facing AI

Third-party independent behavioral safety validation for human-facing AI systems.

Benchmark scope: 79 scenarios | 13 behavioral domains (vulnerability categories).

Send completed PDF to research@ikwe.ai or submit online at ikwe.ai/intake.

C. Model + Stack

Model provider(s) (select all that apply)

OpenAI
Anthropic
Google
Meta / open-source
...

Other model provider

Use system prompts?

Use RAG / knowledge base?

Use tools/actions?

Use fine-tuning?

The Behavioral Safety Layer for Human-Facing AI

Third-party independent behavioral safety validation for human-facing AI systems.

Benchmark scope: 79 scenarios | 13 behavioral domains (vulnerability categories).

Send completed PDF to research@ikwe.ai or submit online at ikwe.ai/intake.

D. Risk / Governance Pressure

What is driving this evaluation? (select all that apply)

Board request
Customer procurement/security review
Partner requirement
Compliance readiness
Incident/near-miss

Other evaluation driver

Any deadlines?

Date or timeline

Who needs to sign off? (select all that apply)

CEO/founder
CTO
Head of Compliance / Risk
Legal
Customer security/procurement

Other sign-off stakeholder

E. Access Feasibility

Preferred testing method

Can you provide sandbox/staging access?

Can you provide 2-3 controlled test accounts?

Are outputs confidential?

The Behavioral Safety Layer for Human-Facing AI

Third-party independent behavioral safety validation for human-facing AI systems.

Benchmark scope: 79 scenarios | 13 behavioral domains (vulnerability categories).

Send completed PDF to research@ikwe.ai or submit online at ikwe.ai/intake.

F. Scope Signals

Approximate scenario volume

Preferred engagement mode

G. Open Text

Briefly describe your system and what you are most concerned about.

List any specific red-lines your organization must enforce.

The Behavioral Safety Layer for Human-Facing AI

Third-party independent behavioral safety validation for human-facing AI systems.

Benchmark scope: 79 scenarios | 13 behavioral domains (vulnerability categories).

Send completed PDF to research@ikwe.ai or submit online at ikwe.ai/intake.

H. Data Handling + Security

Can Ikwe store outputs for analysis?

Can outputs include PII?

Required retention period

Custom retention detail

Compliance constraints (select all that apply)

HIPAA / PHI

PCI

FERPA

GDPR

Minors

No additional constraints

Other compliance constraints

I. Scope + Success Criteria

What does "pass" mean for your team?

Priority behavioral domains (select all that apply)

Anxiety

Depression

Loneliness

Anger

Overwhelm

Grief

Suicidal Ideation

Languages supported