

1-Page Executive Excerpt

EQ Safety Benchmark v2.0 · Proprietary format illustration



• Safety Gate: Conditional Pass

TIER II · CONTROLLED ROLLOUT WITH SAFEGUARDS

Generally safe, but with non-trivial behavioral risks that require attention before scaling to high-sensitivity user populations.

In plain language

- The system handles routine emotional exchanges well, but shows meaningful gaps when users escalate into distress or crisis-adjacent states.
- Crisis routing fails to consistently redirect vulnerable users to appropriate resources — partial stabilization in high-arousal scenarios.
- Mild dependency pressure observed in extended support conversations. Not catastrophic, but warrants monitoring and guardrails before scaling.
- Adversarial inputs destabilize emotional safety behaviors disproportionately. This is the primary remediation priority.

EVALUATION PROFILE

WHERE WE TESTED	HOW WE TESTED	METHOD
Customer support assistant, high-touch consumer context	79 designed conversations · 12 vulnerability domains · 948 scored outputs	EQ Safety Benchmark v2.0 + support-context scenario slice

WHAT EACH TEAM RECEIVES FROM A FULL REPORT

BOARD & EXECUTIVES Clear yes / no / under-conditions call on deployment readiness, backed by a structured audit trail.	SAFETY & LEGAL Evidence that stands up when regulators, partners, or insurers ask how behavioral risk was evaluated before launch.
COMPLIANCE Scenario-level traceability and scoring records for institutional review, audit trails, and policy documentation.	PRODUCT & ENGINEERING Prioritized failure modes with specific scenario IDs — remediation is targeted, not speculative.

Ready for an independent review?

Scope and run an evaluation in 2–3 weeks. Starts with a 30-minute conversation about your deployment context.

[ikwe.ai/audit](#)

Full sample: [ikwe.ai/reports/ikwe-sample-report-public.html](#)

Not shown: client identifiers, full scenario library, response excerpts, raw run data and reviewer annotations.