



Ikwe EQ Safety Evaluation

Illustrative report structure only. No client-identifiable content.



DIMENSION SNAPSHOT

Illustrative score breakdown



Priority findings

1. Crisis routing is inconsistent under higher-intensity prompts

2. Dependency language appears after prolonged reassurance

3. Repair quality is strong once the system recognizes the model's limitations

Public index comparison

Claude reference row	78.3
Illustrative client system	73.0
GPT-4o Study I baseline	51.2
Grok Study I baseline	40.6

Client evaluations stay private but can be positioned against the same public benchmark record.

What happens after scoring

Board

Decision-ready summary

Classification band, main failure modes, and immediate next-step recommendation.

Legal

Due diligence record

Structured evidence package for liability review and internal governance files.

Compliance

Versioned documentation

Artifacts tied to methodology version, scenario set, and review date.

Engineering

Failure mapping

Specific problem patterns translated into remediation priorities and retest needs.

Recommended next steps

1. Add stronger crisis-routing controls before wider rollout.
2. Reinforce user autonomy in emotionally escalated conversations.
3. Harden prompt and policy layers against dependency reinforcement.
4. Re-test after changes are in place before broader deployment.

Boundary note

Not shown in the public sample: exact scenario text, scoring formulas, raw response evidence, client-identifying details, and full remediation workstream notes.

Request Ikwe EQ Safety Evaluation

ikwe.ai/intake • Compare your system against the same benchmark framework.