

# Ikwe EQ Safety Evaluation

Illustrative report structure only & no client-identifiable content.

## Behavioral Safety Validation

This sample shows the structure of an Ikwe EQ Safety Evaluation deliverable: risk classification, dimension scoring, governance framing, and remediation direction.

### Overall Score

73%

Illustrative EQ Safety score

### Safety Gate

Conditional Pass

Issues present, mitigations required

### Tier

Tier II

Launch with safeguards and re-test

## Executive Summary

Stable behavior appears in low-volatility interactions, but escalation pathways show interruption and routing risk. Current profile supports controlled rollout with safeguards, remediation work, and a targeted re-test before broader deployment.

This sample mirrors the report structure clients receive after a completed evaluation cycle: score, risk posture, evidence, next steps.

### Report Profile

Deployment surface

Customer support assistant

Scenario battery

EQSB baseline + deployment slice

Analyst note

Public sample, redacted

## What leadership gets

- Board brief with classification band and threshold language
- Risk scorecard with dimension outcomes and failure mapping
- Evidence pack for procurement, legal, and compliance review
- Remediation roadmap tied to retest readiness

PDF sample • <https://ikwe.ai/reports/ikwe-sample-report-public.pdf>

Intake • <https://ikwe.ai/intake> | Web preview • <https://ikwe.ai/reports/ikwe-sample-report-public.html>

# Illustrative score breakdown



**Priority findings**

- 1. Crisis routing inconsistent under higher-intensity prompts.
- 2. Dependency language appears after prolonged reassurance loop.
- 3. Repair quality is strong once the system recognizes the miss.

| Public index comparison                                                                         |      |
|-------------------------------------------------------------------------------------------------|------|
| Claude reference row                                                                            | 78.3 |
| Illustrative client system                                                                      | 73.0 |
| GPT-4o Study I baseline                                                                         | 51.2 |
| Grok Study I baseline                                                                           | 40.6 |
| Client evaluations stay private but can be positioned against the same public benchmark record. |      |

## What happens after scoring

### Board

#### Decision-ready summary

Classification band, main failure modes, and immediate next-step recommendation.

### Legal

#### Due diligence record

Structured evidence package for liability review and internal governance files.

### Compliance

#### Versioned documentation

Artifacts tied to methodology version, scenario set, and review date.

### Engineering

#### Failure mapping

Specific problem patterns translated into remediation priorities and retest needs.

### Recommended next steps

1. Add stronger crisis-routing controls before wider rollout.
2. Reinforce user autonomy in emotionally escalated conversations.
3. Harden prompt and policy layers against dependency reinforcement.
4. Re-test after changes are in place before broader deployment.

### Boundary note

Not shown in the public sample: exact scenario text, scoring formulas, raw response evidence, client-identifying details, and full remediation workstream notes.

### Request Ikwe EQ Safety Evaluation

[ikwe.ai/intake](https://ikwe.ai/intake) • Compare your system against the same benchmark framework.

<https://ikwe.ai/intake>